

Chapter 11

Approximation Algorithms



Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.

Approximation Algorithms

Q. Suppose I need to solve an NP-hard problem. What should I do?

A. Theory says you're unlikely to find a poly-time algorithm.

Must sacrifice one of three desired features.

- Solve problem to optimality.
- Solve problem in poly-time.
- Solve arbitrary instances of the problem.

ρ -approximation algorithm.

- Guaranteed to run in poly-time.
- Guaranteed to solve arbitrary instance of the problem
- Guaranteed to find solution within ratio ρ of true optimum.

Challenge. Need to prove a solution's value is close to optimum, without even knowing what optimum value is!

11.1 Load Balancing

Load Balancing

Input. m identical machines; n jobs, job j has processing time t_j .

- Job j must run contiguously on one machine.
- A machine can process at most one job at a time.

Def. Let $J(i)$ be the subset of jobs assigned to machine i . The **load** of machine i is $L_i = \sum_{j \in J(i)} t_j$.

Def. The **makespan** is the maximum load on any machine $L = \max_i L_i$.

Load balancing. Assign each job to a machine to minimize makespan.

Load Balancing: List Scheduling

List-scheduling algorithm.

- Consider n jobs in some fixed order.
- Assign job j to machine whose load is smallest so far.



Note: in
Textbook, T_i is
used instead of
 L_i

```

List-Scheduling( $m, n, t_1, t_2, \dots, t_n$ ) {
  for  $i = 1$  to  $m$  {
     $L_i \leftarrow 0$       ← load on machine  $i$ 
     $J(i) \leftarrow \phi$  ← jobs assigned to machine  $i$ 
  }

  for  $j = 1$  to  $n$  {
     $i = \operatorname{argmin}_k L_k$  ← machine  $i$  has smallest load
     $J(i) \leftarrow J(i) \cup \{j\}$  ← assign job  $j$  to machine  $i$ 
     $L_i \leftarrow L_i + t_j$  ← update load of machine  $i$ 
  }
  return  $J(1), \dots, J(m)$ 
}

```

Implementation. $O(n \log m)$ using a priority queue.

5

Load Balancing: List Scheduling Analysis

Theorem. [Graham, 1966] Greedy algorithm is a 2-approximation.

- First worst-case analysis of an approximation algorithm.
- Need to compare resulting solution with optimal makespan L^* .

Lemma 1. The optimal makespan $L^* \geq \max_j t_j$.

Pf. Some machine must process the most time-consuming job. ▪

Lemma 2. The optimal makespan $L^* \geq \frac{1}{m} \sum_j t_j$.

Pf.

- The total processing time is $\sum_j t_j$.
- One of m machines must do at least a $1/m$ fraction of total work. ▪


6

Load Balancing: List Scheduling Analysis

Theorem. Greedy algorithm is a 2-approximation.

Pf. Consider load L_i of bottleneck machine i .

- Let j be last job scheduled on machine i .
- When job j assigned to machine i , i had smallest load. Its load before assignment is $L_i - t_j \Rightarrow L_i - t_j \leq L_k$ for all $1 \leq k \leq m$.



7

Load Balancing: List Scheduling Analysis

Theorem. Greedy algorithm is a 2-approximation.

Pf. Consider load L_i of bottleneck machine i .

- Let j be last job scheduled on machine i .
- When job j assigned to machine i , i had smallest load. Its load before assignment is $L_i - t_j \Rightarrow L_i - t_j \leq L_k$ for all $1 \leq k \leq m$.
- Sum inequalities over all k and divide by m :

$$\begin{aligned}
 L_i - t_j &\leq \frac{1}{m} \sum_k L_k \\
 &= \frac{1}{m} \sum_j t_j \\
 \text{Lemma 1} \rightarrow &\leq L^*
 \end{aligned}$$

Sum of makespan L_k equals sum of all jobs' processing time t_j .

$$\begin{aligned}
 \text{Now } L_i &= \underbrace{(L_i - t_j)}_{\leq L^*} + \underbrace{t_j}_{\leq L^*} \leq 2L^*.
 \end{aligned}$$

\uparrow
 Lemma 2

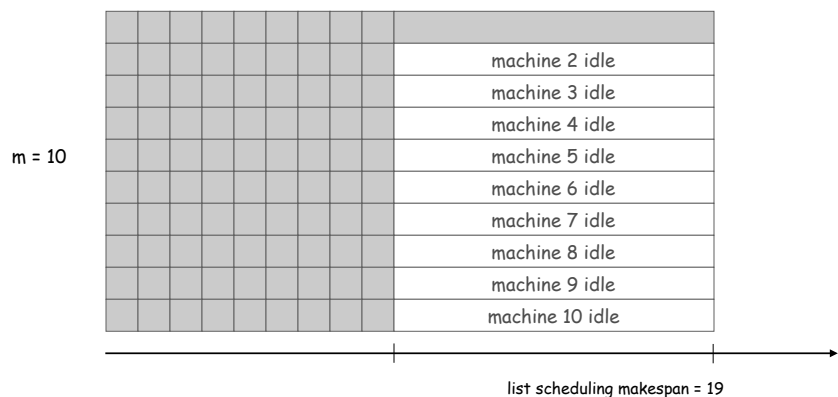
8

Load Balancing: List Scheduling Analysis

Q. Is our analysis tight?

A. Essentially yes.

Ex: m machines, $m(m-1)$ jobs length 1 jobs, one job of length m




9

Load Balancing: List Scheduling Analysis

Q. Is our analysis tight?

A. Essentially yes.

Ex: m machines, $m(m-1)$ jobs length 1 jobs, one job of length m



10

Load Balancing: LPT Rule

Longest processing time (LPT). Sort n jobs in **descending** order of processing time, and then run list scheduling algorithm.

```

LPT-List-Scheduling( $m, n, t_1, t_2, \dots, t_n$ ) {
  Sort jobs so that  $t_1 \geq t_2 \geq \dots \geq t_n$ 

  for  $i = 1$  to  $m$  {
     $L_i \leftarrow 0$       ← load on machine  $i$ 
     $J(i) \leftarrow \phi$   ← jobs assigned to machine  $i$ 
  }

  for  $j = 1$  to  $n$  {
     $i = \operatorname{argmin}_k L_k$     ← machine  $i$  has smallest load
     $J(i) \leftarrow J(i) \cup \{j\}$  ← assign job  $j$  to machine  $i$ 
     $L_i \leftarrow L_i + t_j$  ← update load of machine  $i$ 
  }
  return  $J(1), \dots, J(m)$ 
}

```

11

Load Balancing: LPT Rule

Observation. If at most m jobs, then list-scheduling is optimal.

Pf. Each job put on its own machine. ■

Lemma 3. If there are more than m jobs, $L^* \geq 2t_{m+1}$.

Pf.

- Consider first $m+1$ jobs t_1, \dots, t_{m+1} .
- Since the t_i 's are in descending order, each takes at least t_{m+1} time.
- There are $m+1$ jobs and m machines, so by pigeonhole principle, at least one machine gets two jobs. ■

Theorem. LPT rule is a $3/2$ approximation algorithm.

Pf. Same basic approach as for list scheduling.

$$L_i = \underbrace{(L_i - t_j)}_{\leq L^*} + \underbrace{t_j}_{\leq \frac{1}{2}L^*} \leq \frac{3}{2}L^*. \quad \blacksquare$$

\uparrow
 Lemma 3
 (by observation, can assume number of jobs $> m$)


12

11.2 Center Selection

Center Selection Problem

Input. Set of n sites s_1, \dots, s_n and integer $k > 0$.

Center selection problem. Select k centers C so that maximum distance from a site to nearest center is minimized.



Center Selection Problem

Input. Set of n sites s_1, \dots, s_n and integer $k > 0$.

Center selection problem. Select k centers C so that maximum distance from a site to nearest center is minimized.

Notation.

- $\text{dist}(x, y)$ = distance between x and y .
- $\text{dist}(s_i, C) = \min_{c \in C} \text{dist}(s_i, c)$ = distance from s_i to closest center.
- $r(C) = \max_i \text{dist}(s_i, C)$

Goal. Find set of centers C that minimizes $r(C)$, subject to $|C| = k$.

Distance function properties.


- $\text{dist}(x, x) = 0$ (identity)
- $\text{dist}(x, y) = \text{dist}(y, x)$ (symmetry)
- $\text{dist}(x, y) \leq \text{dist}(x, z) + \text{dist}(z, y)$ (triangle inequality)

16

Center Selection Example

Ex: each site is a point in the plane, a center can be any point in the plane, $\text{dist}(x, y)$ = Euclidean distance.

Remark: search can be infinite!




17

Greedy Algorithm: A False Start

Greedy algorithm. Put the first center at the best possible location for a single center, and then keep adding centers so as to reduce the covering radius each time by as much as possible.

Remark: arbitrarily bad!



18

Center Selection: Greedy Algorithm

Greedy algorithm. Repeatedly choose the next center to be the site **farthest** from any existing center.

```

Greedy-Center-Selection( $k, n, s_1, s_2, \dots, s_n$ ) {
     $C = \emptyset$ 
    Select any site  $s$  and add  $s$  to  $C$ ;
    repeat  $k-1$  times {
        Select a site  $s_i$  with maximum  $\text{dist}(s_i, C)$ 
        Add  $s_i$  to  $C$ 
    }
    return  $C$ 
}

```

↑
site farthest from any center

Observation. Upon termination all centers in C are pairwise at least $r(C)$ apart.

Pf. By construction of algorithm.


19

Center Selection: Analysis of Greedy Algorithm

Theorem. Let C^* be an optimal set of centers. Then $r(C) \leq 2r(C^*)$.

Pf. (by contradiction) Assume $r(C^*) < \frac{1}{2} r(C)$.

- For each site c_i in C , consider ball of radius $\frac{1}{2} r(C)$ around it.
- Exactly one c_i^* in each ball; let c_i be the site paired with c_i^* .
- Consider any site s and its closest center c_i^* in C^* .
- $\text{dist}(s, C) \leq \text{dist}(s, c_i) \leq \text{dist}(s, c_i^*) + \text{dist}(c_i^*, c_i) \leq 2r(C^*)$.
- Thus $r(C) \leq 2r(C^*)$. Δ -inequality $\leq r(C^*)$ since c_i^* is closest center



20

Center Selection

Theorem. Let C^* be an optimal set of centers. Then $r(C) \leq 2r(C^*)$.

Theorem. Greedy algorithm is a 2-approximation for center selection problem.

Remark. Greedy algorithm always places centers at sites, but is still within a factor of 2 of best solution that is allowed to place centers anywhere.

e.g., points in the plane


21

11.4 The Pricing Method: Vertex Cover


Weighted Vertex Cover

Definition. Given a graph $G = (V, E)$, a vertex cover is a set $S \subseteq V$ such that each edge in E has at least one end in S .

Weighted vertex cover. Given a graph G with vertex weights, find a vertex cover of minimum weight.



weight = $2 + 2 + 4$




weight = 11

Pricing Method

Pricing method. Each edge must be covered by some vertex.
Edge $e = (i, j)$ pays price $p_e \geq 0$ to use vertex i and j .

Fairness. Edges incident to vertex i should pay $\leq w_i$ in total.

$$\text{for each vertex } i: \sum_{e=(i,j)} p_e \leq w_i$$



Lemma. For any vertex cover S and any fair prices p_e : $\sum_e p_e \leq w(S)$.

Pf.

$$\sum_{e \in E} p_e \leq \sum_{i \in S} \sum_{e=(i,j)} p_e \leq \sum_{i \in S} w_i = w(S).$$

\uparrow each edge e covered by at least one node in S \uparrow sum fairness inequalities for each node in S

25

Pricing Method

Pricing method. Set prices and find vertex cover simultaneously.

```

Weighted-Vertex-Cover-Approx(G, w) {
  foreach e in E
    pe = 0
    while (∃ edge i-j such that neither i nor j are tight)
      select such an edge e
      increase pe as much as possible until i or j tight
    }
  S ← set of all tight nodes
  return S
}

```

$$\sum_{e=(i,j)} p_e = w_i$$

\downarrow

26

Pricing Method




Figure 11.8

27

Pricing Method: Analysis

Theorem. Pricing method is a 2-approximation.

Pf.

- Algorithm terminates since at least one new node becomes tight after each iteration of while loop.
- Let S = set of all tight nodes upon termination of algorithm. S is a vertex cover: if some edge i - j is uncovered, then neither i nor j is tight. But then while loop would not terminate.
- Let S^* be optimal vertex cover. We show $w(S) \leq 2w(S^*)$.

$$w(S) = \sum_{i \in S} w_i = \sum_{i \in S} \sum_{e=(i,j)} p_e \leq \sum_{i \in V} \sum_{e=(i,j)} p_e = 2 \sum_{e \in E} p_e \leq 2w(S^*). \quad \blacksquare$$

\uparrow all nodes in S are tight \uparrow $S \subseteq V$, prices ≥ 0 \uparrow each edge counted twice \uparrow fairness lemma

28